

Project 2 Design Doc

A Secure File-sharing System

on a Hostile Storage

1. System Design

a. How is a file stored on the server?

Files are stored by users. To first store user info, we encode a User struct into JSON, symmetrically encrypt it with `k_user_encrypt`, pad it until it is divisible by 16 (`AESBlockSize`) according to PKCS #7, HMAC the ciphertext with `k_user_auth`, and store the HMAC appended by the ciphertext on Datastore, at the storage key `ID_user`. To generate those keys, we first deterministically derive `k_password` with PKDF from the password, salted by the unique username s.t. different users have different keys even if they share the same passwords. We then deterministically derive the other keys with HKDF from `k_password`, with the purpose array set as the encoded version of string 'user_encrypt' for `k_user_encrypt`, that of 'user_auth' for `k_user_auth`, and 'user_storage' for `k_user_storage`. `ID_user` is the UUID of the HMAC of the (bytes of) username with key `k_user_storage`.

We also generate a PKE keypair, `k_pub` and `K_private`, and a DS keypair, `k_DS_pub` and `K_DS_private`, for each user for sharing. We publish `k_pub` on Keystore at the storage key `k_pubkey`, which is an UUID generated from the first 16B of the hash of the unique username appended by string `'public_key'`. We publish `k_DS_pub` at the storage key `k_DSkey`, which is another UUID generated from the first 16B of the hash of the unique username appended by string `'DS_key'`. We store `K_private` and `K_DS_private` in the secured and authenticated User struct.

For a user to retrieve their User struct, they enter their username and password to generate `ID_user` through the aforementioned deterministic process. If `ID_user` exists and the password is correct, the `user_struct` will be decrypted and depadded (according to PKCS #7) upon retrieval.

To store a file, a user calls `User.StoreFile()` with their User struct. The file is encoded into a byte array and splitted, each volume encrypted and then its ciphertext authenticated by its MAC. The split is in blocks of 1GB to ease future appending. The last block is padded to 1GB with the length of the pad in bytes (following PKCS #7). The encryption is through AES-CBC encryption, with the 16B IV generated by a RNG for every volume and the key `k_volume` deterministically derived through HKDF from `k_file`, with the purpose array set as the encoded version of string `'volume_encryption'` appended by the index of the volume. Therefore, we have different IVs and keys for all volumes. `k_file` is a random 128b (16B) symmetric key from our RNG. The HMAC key is `k_volume_MAC`, again deterministically derived through HKDF from `k_file`, with the purpose array set as the encoded version of string `'volume_authentication'` appended by

the index of the volume. We define a Volume struct that contains the ciphertext and the HMAC as two arrays, and the number of padded bytes as uint32.

We fetch `k_pub` and `k_DS_pub` for the user, append `k_file` to 16B of random padding in front to prevent IND-CPA, PKE encrypt the 32B `k_file_front_padded` with `k_pub` to output `pke_k_file`, and Digitally Sign `pke_k_file` with `K_DS_private` to output `ds_k_file`. We store `ds_k_file` and `pke_k_file` in a SignedKey struct, and store it in a marshalled map of SignedKeys on Datastore at `ID_k`, where `ID_k` is a random UUID. The key in the map is the username. We store `ID_k` in map `AES_key_storage_keys` in the User struct, with both keys being the filename, and update the User struct on Datastore.

The encrypted volumes are then put together as an array of Volume structs, marshalled again and stored on Datastore. The storage key is `ID_file`, which is generated from the hash of the string representation of `ID_k` converted to a byte array.

If a user calls `StoreFile` on a filename that already exists, the file content is overwritten.

b. How does a file get shared with another user?

The sharer uses the `ShareFile` method, which will generate an access token to be sent to the sharee and is called whenever the owner of the file selects a file to be shared. The token only permits the sharee to read, write and share the file, and nobody else.

We first retrieve the master key, `k_file`. We find the storage key `ID_k` of all encrypted AES keys (`k_file`'s) for the file in a map in the owner's User struct, find the owner's SignedKey with their username, and verify and decrypt it with the owner's public DS key and private PKE key.

We then PKE encrypt k_{file} with the recipient's public key k_{pub} generated when the user was initiated, Digitally Sign the ciphertext with the sharer's $k_{DS_private}$, create a SignedKey struct with the signature and the ciphertext, and add it to the SignedKey map at ID_k , with the recipient's username as the key. We then marshal and store the updated SignedKey array at ID_k . The token is a SignedKey struct containing two variables: ID_k encrypted by the recipient's k_{pub} , and another version that is Digitally Signed with the sharer's $k_{DS_private}$. Lastly, if the sharer owns the file, we note the direct recipient of that file in a map in the owner's User struct.

When the recipient receives the token using ReceiveFile, they will first come up with their own filename for it. If they already have a file with that name, err. If not, they verify the signature with the sharer's k_{DS_pub} and decrypt the message using their own private key $K_{private}$. If verification works and the file exists, a mapping from an arbitrary new filename for the recipient to ID_k is added to their User struct. If verification fails or the file does not exist, return an error.

c. What is the process of revoking a user's access to a file?

The owner will select the user whose access will be removed and then call RevokeFile. This method removes the ability to read, write and share for the selected user.

It first checks whether the user owns the file. If they do, it then generates a new 16B symmetric key k_{file} , and gets the segmented plaintext like in LoadFile. encrypts and authenticates the plaintext with it like in the storing process, encrypts k_{file} like in regular sharing with every direct sharee's public keys k_{pub} except for the revokee's

(according to a map in the owner's secure User struct), signs all those ciphertexts, and stores the updated SignedKey array back at ID_k.

Returns an error if the selected user does not own the file or if file does not exist.

d. How does your design support efficient file append?

Since the file will be segmented into volumes of 1GB in size, we won't have to decrypt or authenticate the entire file. Instead, we only need to decrypt the last 1GB of data, append to that data, then re-encrypt it. This will be faster than if we had to decrypt the entire file.

2. Security Analysis

a. Attack on Storage

If user A stores a file, user M might store their own file to have the same name as the original file's. Our design prevents problems by not incorporating filenames in storage locations etc, and only stores it in a secure User struct, which is part of the Trusted Computing Base.

b. Attack on Sharing

The attacker has the ability to read the token in transit. Thus, they can use it to gain unauthorized access to shared files by . In order to prevent this, our design uses the public key of the recipient (stored in Datastore) to encrypt it so it can only be decrypted by the user to which the token is intentionally sent.

c. Attack on Revocation

User A could share a file with user B then revoke access. User B could then try a denial of service attack by arbitrarily changing contents of the file. To detect that, we have designed HMACs for each volume of the file, where the keys are derived from k_{file} , a secure master key whose possession is equivalent to legal access. It is an unpredictable random number and is stored elsewhere on Datastore, encrypted by RSA. Therefore, M would not create valid MACs for the modified file, and A will detect any change after an unauthorised modification and MAC authentication fails.